

AWS Solutions Architect Associate

Practice Questions

Question 1

IAM role in the company account, cross-account trust with partner Your company has developed a product that involves secure data transfer to individual partners. Each partner has unique, sensitive data, which is stored in dedicated, isolated S3 buckets to ensure data segregation and security. Considering this setup, which permissions implementation would best minimize the risk of any partner accessing inappropriate data that does not belong to them?

- A. S3 bucket policy applied to each bucket with cross-account permissions
- B. S3 bucket ACL + object ACLs with cross-account permissions
- C. IAM user in the company account, credentials delivered to partner
- D. IAM role in the company account, cross-account trust with partner

Explanations

A: The S3 bucket policy is unique to each bucket, and can be customized to allow least-privilege access to a single partner without the possibility of the partner accessing any other S3 bucket.

B: Like S3 bucket policies, bucket ACLs and object ACLs are specific to the scope they're configured to, but do not allow for least-privilege design, as the grantee can only be an entire 12-digit AWS account ID.

C: This solution, while it can be least-privileged, allows for the possibility of "extra" permissions being associated with the IAM user.

D: This solution is similar to C in that it is possible for least privilege, but also possible for extra or unintended permissions.

Question 2

AWS Trusted Advisor A company is using AWS to host its multi-tier web application. The application architecture includes Amazon EC2, Amazon RDS, and Amazon S3. The company has a strict compliance policy that requires all resources to adhere to specific standards. They are looking for a solution that can continuously monitor and evaluate the configurations of their AWS resources against the desired configurations and provide alerts when non-compliance is detected. Which AWS service should the security team use to meet this requirement?

- A. Amazon Inspector
- B. AWS Config
- C. Amazon CloudWatch

D. AWS Trusted Advisor

Explanations

A: Amazon Inspector is primarily used for security assessments and vulnerability management of EC2 instances and applications running on those instances. It does not provide continuous monitoring and evaluation of resource configurations against compliance standards.

B: Config is designed to assess, audit, and evaluate the configurations of your AWS resources. AWS Config continuously monitors and records your AWS resource configurations and allows you to automate the evaluation of recorded configurations against desired configurations. This service can alert you whenever non-compliant resources are detected, making it the ideal solution for the company's requirement.

C: While CloudWatch provides monitoring and operational data for AWS resources, it does not evaluate resource configurations against compliance standards. CloudWatch is more focused on performance monitoring and operational health.

D: Trusted Advisor provides recommendations to help you follow AWS best practices. While Trusted Advisor does offer some checks related to security and compliance, it does not provide the continuous monitoring and evaluation of specific resource configurations against compliance standards like AWS Config does.

Question 3

An application is deployed to EC2 instances in a private VPC subnet in the us-west-2 region. There is a functional requirement to access S3 buckets deployed into the ap-southeast-1 region. There is a security requirement for the end to end traffic to remain private. Which combination of steps can meet the functional and security requirements? (pick three)

- A. Deploy a VPC Gateway endpoint into the us-west-2 VPC
- B. Deploy a VPC in the ap-southeast-1 region with an S3 Interface endpoint
- C. Configure a route table entry in the us-west-2 VPC to direct S3 traffic to the Gateway endpoint
- D. Configure a VPC peering connection between the us-west-2 and ap-southeast-1 VPCs
- E. Configure the application to use the ap-southeast-1 region when accessing S3 buckets
- F. Configure the application to use the ap-southeast-1 Interface endpoint DNS when accessing S3 buckets

Explanations

A, C, E: This may seem like a functional solution that meets the security requirements. Unfortunately, while Gateway endpoints do proxy S3 traffic to the service API endpoint (meeting the security requirement), they cannot accept traffic for buckets deployed in any region outside that of the endpoint.

B, D, F: This solution is more complicated. S3 endpoints, whether they are Gateway or

Interface, cannot route traffic to cross-region buckets. Furthermore, Gateway endpoints cannot accept traffic from a remote VPC via a peering connection, so the Interface endpoint is the only functional option. This creates an ENI to accept the S3 traffic, with an associated DNS entry which can be used from the remote VPC. The traffic remains private, meeting the security requirement.

Question 4

A financial services company is planning to deploy a new application on AWS. The application will handle sensitive customer data, including financial transactions. The company's security team wants to ensure that the application and its environment adhere to strict security best practices from the moment they are deployed. They are particularly interested in minimizing human error and ensuring consistent application of security controls across their AWS environment. Which AWS service should the security team use?

- A. AWS Security Hub
- B. AWS Config
- C. AWS Systems Manager
- D. AWS CloudFormation

Explanations

A: AWS Security Hub provides a comprehensive view of your security state within AWS and can help you check your environment against security industry standards and best practices. However, it does not automate the deployment of security controls or resources.

B: While AWS Config can be used to assess, audit, and evaluate the configurations of your AWS resources, it is more focused on monitoring and recording the configurations of AWS resources and evaluating those configurations for compliance with desired guidelines. It does not automate the deployment of resources or security controls.

C: Although AWS Systems Manager provides visibility and control of your infrastructure on AWS, it is primarily used for operational tasks such as patch management, application deployment, and resource configuration. It does not specifically automate the deployment of security controls based on predefined security standards.

D: CloudFormation allows you to model and set up your resources so that you can spend less time managing those resources and more time focusing on your applications. You can use AWS CloudFormation to automate the deployment of resources in a secure configuration, ensuring that all resources comply with the company's security standards.

Question 5

Your company securely stores sensitive data in S3. There is a strict requirement to ensure that all data is encrypted at rest using the method which allows for the most customer control of the encryption. Which of the following encryption configurations would best meet the requirement for safeguarding this sensitive information against unauthorized access?

- A. SSE-S3
- B. SSE-KMS (AWS Managed)
- C. SSE-KMS (Customer Managed)
- D. Client-side encryption

Explanations

A: SSE-S3 uses AES-256 with AWS owning the entire chain of trust (Root CA, master key, data key). The encryption is performed in the S3 service. This is slightly better than no encryption at all.

B: This is a better option than A. AWS owns the Root CA, and the customer owns the remainder of the chain of trust. The encryption is performed in the S3 service. The customer does not have the ability to customize the key rotation schedule.

C: This is more secure than B. AWS owns the Root CA, and the customer owns the remainder of the chain of trust. The encryption is performed in the S3 service. The customer DOES have the ability to customize the key rotation schedule, or to even import key material from an external Key Management Infrastructure.

D: This solution does not use AWS for encryption whatsoever, and, like C, allows the customer to own the entire chain of trust. This option requires the customer to directly encrypt the data, instead of it being encrypted on the server side, and thus provides more customer control.

Question 6

An on-premises application stores data containing PII on an EFS file system, reachable through a site-to-site VPN. There is a requirement for each file to have certain fields tokenized as quickly as possible to avoid the raw data being easily accessed. The tokenization must be completed entirely within AWS. Which of the following workflows would meet the requirements?

- A. Deploy a Lambda function which accesses EFS directly and uses KMS to encrypt the fields inline in each file.
- B. Deploy a Lambda function which accesses EFS directly. Generate a random string to replace each field, then extract the plaintext from each field. Encrypt the field using KMS and insert into a DynamoDB table with the random string as the partition key.
- C. Use an EMR cluster to access the EFS filesystem and use KMS to encrypt the fields inline in each file.
- D. Deploy an AWS Batch job which accesses EFS directly. Generate a random string to replace each field, then extract the plaintext from each field. Encrypt the field using KMS and insert into a DynamoDB table with the random string as the partition key.

Explanations

A: This is functional as far as AWS is concerned, and can be executed in near real-time. It doesn't accomplish tokenization, which is a form of pseudonymization where certain

sensitive data is separated from the rest of the data to make it no longer identifiable.

B: Like A, this is functional and near real-time, but it also accomplishes the pseudonymization by separating some data into DynamoDB with a link back to the text string in the original data.

C: Similar to A, this is functional but does not meet either the requirement for tokenization or the requirement for "as soon as possible".

D: While this accomplishes tokenization, it is not done in near real-time fashion.

Question 7

AWS has emailed an abuse notice to your root account identifying a compromised EC2 instance. Your security team wants to prevent the instance from performing any further activity while not affecting the running OS or other workloads. Which of the following actions would meet the requirement?

- A. Create a security group with no inbound or outbound rules. Replace the compromised instance's security group with the newly created one
- B. Create a security group with no inbound or outbound rules. Attach the newly created security group to the compromised instance
- C. Modify the subnet network ACL and block all outbound traffic
- D. Modify the subnet network ACL and block all inbound and outbound traffic

Explanations

A: A lack of Security group rules can block all traffic to and from the instance, effectively isolating it from the rest of the VPC network.

B: This sounds like a functional solution, but the effective Security group rules are the union of all rules, then applied by specificity. This task accomplishes nothing.

C: This is a functional mitigation step. However, it will affect all other instances in the subnet, while still allowing the compromised instance to explore within the subnet.

D: This solution is effectively identical to C in that it partially stops the instance from communicating, but also impacts other resources.

Question 8

You've been asked to design the monitoring of an internal application which runs on an EC2 instance in a private subnet, listening on port 8080. This monitoring will be used to calculate availability of the application and be used for business continuity purposes as well as automated failover. Which of the following recommendations will meet the requirement in the most reliable manner? (pick two)

- A. Write a script to run via cron every minute which checks the listener port and uploads as a custom CloudWatch metric
- B. Configure the CloudWatch agent to push a metric which tests the listener
- C. Provision a CloudWatch Synthetics canary to test the listener port

- D. Configure a CloudWatch alarm from the appropriate metric. Configure a Route 53 Health Check with the alarm as the source
- E. Configure a Route 53 Health Check with the CloudWatch metric as the source

Explanations

A: This is a plausible choice for delivering the listener uptime metrics to CloudWatch. There is some question whether a shell script via cron is the most reliable method for delivery.

B: CloudWatch agent configuration does not support custom metrics of this type.

C: This is also a functional choice for delivering uptime metrics, and is likely more reliable than the cron job option described in A.

D: This assumes the metric has been delivered to CloudWatch. Once an alarm has been created, it can be used as a source for a Route 53 Health Check.

E: Route 53 Health Checks can use CloudWatch alarms as a source, but not a CloudWatch metric directly.

Question 9

An application has a requirement for cross-region recovery of a DynamoDB table in case of disaster. The requirements include being able to recover any deleted items. The table in the source region has DynamoDB Streams enabled. What recommendation should be made to meet the requirement?

- A. Configure the global table option and add the disaster recovery region
- B. Provision a Lambda function triggered by the DynamoDB stream. Parse the stream entries and propagate Create and Update operations only to the remote table.
- C. Provision a Lambda function triggered by the DynamoDB stream. Parse the stream entries and propagate all operations to the remote table.
- D. Provision an EventBridge rule to capture all table events. Forward the events to a Lambda function which parses the events and applies them to the remote table.

Explanations

A: This would be a perfect solution, except that the global table feature also propagates item deletes.

B: This recommendation is functional, and accounts for the requirement to recover deleted items.

C: This solution is mostly functional, but similar to A in that item deletes would also be propagated.

D: This solution is yet another way to propagate all changes, similar to A and C, but will not meet the requirement for recovering deleted items.

Question 10

A stateful eCommerce EC2 application has been deployed using Auto Scaling behind an Application Load Balancer. During scaling, customers are complaining about their shopping carts suddenly being empty. What steps could be taken to ensure shopping cart data is retained during EC2 scaling events?

- A. Change the Auto Scaling group to a steady state group (same minimum, maximum, desired instances) to match the peak traffic load
- B. Enable duration-based stickiness in the EC2 Target group to ensure all requests from the same client are routed to the same target
- C. Enable application-based stickiness in the EC2 Target group to ensure all requests from the same client are routed to the same target
- D. Migrate the shopping cart data to a DynamoDB table and reference the table instead of the local application cache

Explanations

A: This solution can be an incremental improvement, but the true fix for "scaling causes problems" shouldn't be "don't scale at all", and this does not account for AWS-related events that cause instance replacement.

B: Enabling TG stickiness can improve the situation by ensuring that each client stays on the same target for the duration set, but does not account for scaling-in activity.

C: Similar to B, enabling TG stickiness can improve the situation by ensuring that each client stays on the same target for the duration set, but does not account for scaling-in activity. This could be slightly better than B if the cookie is maintained for a long period of time.

D: None of the other solutions truly address the root cause of the problem, which is that the application is stateful. Re-architecting the application to be stateless by moving the shopping cart persistence to a separate, reliable/durable storage service will result in the best overall solution from a resilience perspective (but could have tradeoff implications on performance due to extra latency of accessing shopping cart data).

Question 11

A web application is being architected for deployment into AWS. There is a requirement to implement throttling of web requests when the back end latency increases beyond a specific threshold, in order to ensure an optimal user experience. Which of the below recommendations can meet this requirement?

- A. Application Load Balancer listener rules
- B. CloudFront distribution behaviors
- C. WAF Web ACL rate-based rule
- D. API Gateway usage plans
- E. A custom solution is required

Explanations

A: Listener rules have several options, including static responses, redirects or forwarding to a Target Group, but no option for rejecting requests outright.

B: While CloudFront distributions can be associated with a WAF, they have no intrinsic filtering capability beyond geo blocking.

C: This sounds like an acceptable solution, except that the throttling must be configured via a static threshold, which will not account for the busy-ness of the back end systems.

D: API Gateway usage plans are similar to WAF rate-based rules in that a static threshold must be set and anything beyond that is throttled. It cannot apply a dynamic rule based on an existing metric.

E

Question 12

To improve reliability, you've recommended deploying an application using immutable infrastructure. Which of the following is NOT a benefit of deploying application infrastructure with immutability?

- A. Increased security
- B. Simplified deployments
- C. Increased performance
- D. Increased scalability

Explanations

A: Enforcing replacement of infrastructure during every deployment limits the risk of compromise.

B: Deployments become simpler when there are no inline updates to be performed or verified.

C: While immutability has many advantages, there are no guarantees of improved performance.

D: It is easier to scale by replacing infrastructure than to attempt inline updates, especially from a QA perspective.

Question 13

As part of a new AWS application architecture, you've been asked to recommend a solution that can balance the needs of two requirements: 1. SQL Server Database software with minimum operational overhead 2. Access to the OS and software for regulatory purposes Which recommendation will meet both requirements?

- A. RDS with the SQL Server Database engine selected
- B. RDS with the Custom Engine selected
- C. EC2 with a manual installation of SQL Server Database
- D. AWS does not have any service or feature that can meet both requirements

Explanations

A: The RDS choice of the SQL Server Database engine does not grant access to the underlying Operating System, so this would not meet the requirements.

B: The Custom Engine option for RDS allows customers to access the underlying OS for SQL Server Database, while providing many of the benefits of RDS as a managed service.

C: While this is a functional solution, it does not address the requirement of low operational overhead.

D: The RDS Custom Engine feature can indeed meet both requirements, and this is why it is important to stay current with AWS service and feature updates.

Question 14

An IoT application is deployed using AWS IoT Core, using the MQTT protocol. There is a requirement to provide a single, static IP address, as DNS will not be used. The solution must emphasize high performance. Which combination of steps will meet the requirements? (pick three)

- A. Deploy an IoT Interface endpoint in a VPC within the same region as the IoT Core resource
- B. Deploy an NLB in the VPC using the IoT Interface endpoint in a target group
- C. Deploy an ALB in the VPC using the IoT Interface endpoint in a target group
- D. Configure a Route 53 Alias record for the IoT Interface endpoint
- E. Deploy a CloudFront distribution with the ALB endpoint in an origin group
- F. Deploy a Global Accelerator with the NLB IP address in an origin group

Explanations

A: VPC Interface endpoints do indeed support the IoT Core resources, so this is a possible step.

B: This step assumes that A was already performed, and is also a possibility, as NLB->Interface is a primary implementation pattern for PrivateLink.

C: This step also assumes that A was completed, and is also a supported configuration.

D: This step ignores the requirement to use a static IP address, and therefore does not work as a solution.

E: This step is functionally possible but ignores the static IP address requirement.

F: The Global Accelerator provisions a static AnyCast IP for the endpoint, which can indeed be used in the client IoT application configuration.

Question 15

A workflow requires a Lambda function to be invoked as a target of an EventBridge rule, and will rarely be invoked more than once at the same time. The Lambda function must execute as quickly as possible as a highest priority, but also must optimize for cost. What recommendations should be made to meet these requirements? (pick three)

- A. Configure the Lambda function with more memory
- B. Configure the Lambda function with less memory
- C. Configure the Lambda function with Provisioned Concurrency at a low number
- D. Configure the Lambda function with a Provisioned Concurrency at a high number
- E. Configure the Lambda function with a timeout at the maximum execution time
- F. Configure the Lambda function with a timeout at 95th percentile of maximum

Explanations

A: More memory = more vCPU, which improves performance, meeting the top priority.

B: This optimizes for cost, but not for performance.

C: Provisioned Concurrency ensures there is a prewarmed copy of the function ready to execute, which reduces the startup latency, thus decreasing execution time and improving performance. A low number here would match the expected performance profile.

D: Like C, this improves performance, but since the function is not executed concurrently often, this is a cost risk.

E: This would optimize the reliability of the function, but would end up costing more for failed function executions.

F: This is generally a good recommendation to optimize for cost by assuming that the high runtime outliers are likely in an error condition, so why pay for that extra execution time?

Question 16

An application has a requirement for high storage IOPS, in order to accommodate parallel reads and writes of many small data files (approximately 4Kb each). Which storage solution would NOT be appropriate for a direct performance test of the workload?

- A. FSx for Lustre provisioned with default configuration
- B. EBS Provisioned IOPS io2 Block Express volume
- C. EC2 instance provisioned with an NVMe SSD Instance Store volume
- D. EFS provisioned with General Purpose performance mode

Explanations

A: FSx is widely regarded as high-performing with low latency. However, the default configuration will place all files below 1GiB in the same stripe, reducing the overall performance capacity.

B: The io2 Block Express volume is the next generation of EBS and supports up to 256,000 provisioned IOPS at 16k write blocks.

C: NVMe SSD is very high-performing for IOPS (at 4k write blocks, matching the data size) and throughput, but dependent on the size of the EC2 instance. There will be no network-related latency for this storage option.

D: General Purpose performance mode actually has lower per-operation latency than Max I/O, which is why it was selected as a candidate for the test.

Question 17

A gaming company is developing a high-performance, real-time multiplayer game. The game requires low latency and high throughput to ensure a seamless experience for players worldwide. The development team is focused on optimizing the game's architecture. The team needs to select an AWS service that allows them to deploy their game servers in a way that maximizes performance for players, while accounting for the worldwide distribution of the player base. Which AWS service should the development team use?

- A. Amazon EC2 Auto Scaling
- B. Amazon CloudFront
- C. AWS Global Accelerator
- D. Amazon RDS

Explanations

A: While Auto Scaling allows applications to automatically adjust capacity to maintain steady, predictable performance at the lowest possible cost, it primarily focuses on scaling resources rather than optimizing for low latency and high throughput across geographically distributed users.

B: CloudFront is a content delivery network (CDN) service that securely delivers data, videos, applications, and APIs to customers globally with low latency and high transfer speeds. While it is useful for delivering static and dynamic content, it is not specifically designed for optimizing game server performance and player latency in real-time multiplayer games.

C: The Global Accelerator improves the availability and performance of applications with local or global users by directing traffic to optimal endpoints over the AWS global network. It leverages the AWS backbone network to route traffic efficiently and reduces latency by providing static IP addresses that act as a fixed entry point to application endpoints in one or more AWS Regions. This service is ideal for real-time multiplayer games that require low latency and high throughput, aligning with the concept of mechanical sympathy.

D: RDS is a managed relational database service that makes it easy to set up, operate, and scale a relational database in the cloud. It provides cost-efficient and resizable capacity while automating time-consuming administration tasks such as hardware provisioning, database setup, patching, and backups. While important for data storage and management, Amazon RDS does not directly contribute to minimizing latency or maximizing throughput for real-time multiplayer games.

Question 18

A company has the following AWS migration goals for compute resources: 1. Evolve technology when appropriate 2. Start with familiar technology and learn from there 3. Reduce KTLO (Keep The Lights On) activities over time 4. Agility is top priority, followed by

cost optimization Which of the following compute cost strategies would be appropriate for the company?

- A. Purchase 1 year EC2 reservations
- B. Purchase 3 year EC2 reservations
- C. Purchase 1 year EC2 savings plans
- D. Purchase 1 year Compute savings plans

Explanations

A: 1 year EC2 reservations will help reduce cost, but are limited to a single EC2 instance type.

B: 3 year EC2 reservations provide a larger discount than 1 year, but lock you into a single instance type with some flexibility to change within the instance family, but not much.

C: A savings plan is less discount than a reservation, but more flexible as it covers an entire instance family automatically.

D: A Compute savings plan covers all EC2, Fargate and Lambda, and while the discount is less than EC2 savings plans, it absolutely allows the organization to be more agile in the choice of compute services.

Question 19

When designing a multi-tier application, all non-production environments must be architected to optimize for cost after hours and on weekends. The application tiers include a web proxy, application runtime, and a database. Which of the following recommendations will balance cost optimization with low operational effort? (pick three)

- A. Application Load Balancer
- B. EC2 with nginx using an Auto Scaling group
- C. Application deployed on EC2 instance
- D. Application deployed on EC2 using an Auto Scaling group
- E. RDS database instance
- F. Aurora Serverless

Explanations

A: The ALB, while it is a managed service, does have a minimum footprint on cost even with zero traffic, but the operational overhead of managing this solution is also low/zero.

B: It may seem counterintuitive, but if there is no traffic, the EC2-based web proxy solution can scale better to 0 total cost, by setting the number of desired instances to 0. The overhead of managing this solution is relatively high.

C: Deploying the application manually using a single EC2 instance is not going to scale to 0 well, if at all.

D: An Auto Scaling group (similar to B) can scale to zero with the correct configuration.

E: RDS has no intrinsic features for scaling to 0. It is possible to stop a database instance, but

only for a 1 week period, and not automatically.

F: Aurora Serverless can be configured to reduce the number of ACUs to 0 during periods of idle activity, and entirely automated, so no operational overhead involved.

Question 20

One of your memory-bound applications needs to have a performance baseline established. The application runs on EC2 instances, and you've been asked to design performance baseline tests that include memory usage metrics from the EC2 Operating System. The memory usage metrics will then be used to recommend the appropriate instance size to ensure cost efficient resource usage. How can you ensure the memory metrics are collected?

- A. Use the default CloudWatch metrics for EC2
- B. Enable detailed monitoring on the EC2 instance
- C. Install/configure the CloudWatch Agent on the EC2 instance
- D. Install/configure the CloudWatch Logs Agent on the EC2 instance

Explanations

A: The default CloudWatch metrics are pushed from the hypervisor, and do not include any metrics from inside the guest OS, such as memory usage, or cpu usage breakdown into individual metrics.

B: Detailed monitoring changes the default CloudWatch period from 5 minutes to 1 minute, but does not include any new metrics.

C: The CloudWatch Agent has the capability to deliver memory usage metrics to CloudWatch with a configurable polling period.

D: The CloudWatch Logs Agent is a previous-generation utility for streaming OS or application logs to CloudWatch, and would not be appropriate for memory usage metric delivery.

Question 21

After migrating an on-premises application to EC2, a company needs to identify the method for accessing the instance operating systems as the fleet scales. The method must also scale with the number of instances and optimize for operational complexity. Which recommendation would be appropriate to meet the requirement?

- A. Use Security Group rules and SSH
- B. Use EC2 Serial Console
- C. Use Systems Manager Run Command
- D. Use Systems Manager Session Manager

Explanations

A: This is functionally equivalent to the on-premises infrastructure, which also does not scale as it requires network connectivity to all instances.

B: The EC2 Serial Console is only usable on a single instance, acting as a console connection. This may be useful for a non-bootable instance, but does not scale.

C: Run Command can be deployed against an arbitrary number of instances whether they are on-premises or EC2. It also scales with the number of instances.

D: Session Manager can replace SSH, without the same network connectivity requirements, which scales better, but it is not functionally scalable.

Question 22

Your company's application is using several custom AWS tags that your accounting department would like to generate reports from. The application runs in several AWS accounts that are all part of the same AWS Organization. What action can ensure that the AWS tags are available as part of the AWS monthly bill?

- A. Enable User-defined Cost Allocation Tags for each of the AWS tags in the AWS Organizations Management account
- B. Enable User-defined Cost Allocation Tags for each of the AWS tags in all of the AWS Organizations member accounts
- C. Enable AWS-managed Cost Allocation Tags for each of the AWS tags in the AWS Organizations Management account
- D. No action is required, all tags are automatically made available to the Cost Explorer tool and Cost and Usage Reports

Explanations

A: AWS tags are not automatically converted into Cost Allocation Tags, and so this task must be done manually, which makes this task not only appropriate, but required.

B: Similar to A, the task must be performed manually, but it is not required for all member accounts, just the Management account.

C: The custom tags are not available as AWS-managed Cost Allocation Tags, and so this solution would not be plausible.

D: Tags must be enabled as Cost Allocation Tags manually, and so this solution will not work.